

# U.S. Census Data on Education, Finance, and Jobs

Team-20 Members:

Yan Zheng (zheng23001)

Fahimeh Gholami (fahimeh22000)

Yilin Lai (yilin23000)

## Introduction

---

In this project, a large dataset from the U.S. Census was analyzed. The dataset contains information about **education**, **jobs**, and **income** over many years. The connection between education and income was explored, which can assist businesses and governments in making better decisions. As the dataset is very large, **Apache Spark** was used to process it quickly, and **MongoDB** was used to store the data efficiently. **Python** was used for data analysis and creating predictions. The project is significant because it involves a large amount of data that needs to be processed quickly. Spark and MongoDB were used to handle and analyze the data, enabling useful patterns to be identified.

## Methodology

---

### Dataset Details

Data from the U.S. Census was used, which contains:

- **Education:** Data about the highest level of education attained.
- **Finance:** Information about income levels.
- **Industry:** Data about job types.

The dataset contains millions of records across several years, requiring Big Data tools to process effectively.

### Solution Architecture

A system was built using **Apache Spark** and **MongoDB**:

- **Apache Spark:** Spark was used to process the data quickly and efficiently.
- **MongoDB:** MongoDB was chosen to store the data due to its ability to manage large volumes of data.

To further improve the speed of processing:

- **Partitioning** was applied to the data by **year**, which sped up the queries.
- **Indexes** were created for fields like **Year** and **Education Level** to make data searches faster.

# Implementation

---

## Step-by-Step Execution

### 1.Data Ingestion:

The data was cleaned and organized and then stored in MongoDB.

```
# -----  
# Part 1: Data Ingestion via Pandas and MongoDB (using PyMongo)  
# -----  
import pandas as pd  
from pymongo import MongoClient  
  
# Read CSV files  
finance_df_pd = pd.read_csv('./Finance.csv')  
industry_df_pd = pd.read_csv('./Industry.csv')  
education_df_pd = pd.read_csv('./Educationv.csv')  
  
# Check for missing values  
print("Finance Missing Values:")  
print(finance_df_pd.isnull().sum())  
print("Industry Missing Values:")  
print(industry_df_pd.isnull().sum())  
print("Education Missing Values:")  
print(education_df_pd.isnull().sum())  
  
# Connect to MongoDB and select the target database and collections  
client = MongoClient("mongodb://localhost:27017/")  
db = client["regional_economy"]  
finance_collection = db["finance"]  
industry_collection = db["industry"]  
education_collection = db["education"]  
  
# Delete existing documents in collections  
finance_collection.delete_many({})  
industry_collection.delete_many({})  
education_collection.delete_many({})  
  
# Insert data into MongoDB collections  
finance_data = finance_df_pd.to_dict("records")  
finance_collection.insert_many(finance_data)  
  
industry_data = industry_df_pd.to_dict("records")  
industry_collection.insert_many(industry_data)  
  
education_data = education_df_pd.to_dict("records")  
education_collection.insert_many(education_data)  
  
# Create indexes on frequently used fields (e.g., Year)  
finance_collection.create_index([("Year", 1)])  
industry_collection.create_index([("Year", 1)])  
education_collection.create_index([("Year", 1)])
```

## Using Apache Spark

---

After the data was inserted into MongoDB, **Apache Spark** was used to load and process the data. Spark transformations and queries were performed on the data to explore insights such as trends over time and income analysis.

```
# -----  
# Part 2: Spark Application with MongoDB Integration & Performance Optimization  
# -----  
from pyspark.sql import SparkSession  
from pyspark.sql.functions import col, broadcast  
import matplotlib.pyplot as plt  
import time  
  
# Create SparkSession with MongoDB connector configurations  
spark = SparkSession.builder \  
    .appName("RegionalEconomyAnalysis") \  
    .config("spark.mongodb.input.uri", "mongodb://127.0.0.1/regional_economy.finance") \  
    .config("spark.mongodb.output.uri", "mongodb://127.0.0.1/regional_economy.finance") \  
    .config("spark.jars.packages", "org.mongodb.spark:mongo-spark-connector_2.12:2.4.2,org.mongodb:mongodb-driver") \  
    .getOrCreate()  
  
# Adjust the shuffle partitions; reducing from default 200 to 50 for better performance  
spark.conf.set("spark.sql.shuffle.partitions", "50")  
  
# -----  
# Read MongoDB Data (original functionality maintained)  
# -----  
finance_df = spark.read.format("mongo") \  
    .option("collection", "finance") \  
    .load()  
  
industry_df = spark.read.format("mongo") \  
    .option("uri", "mongodb://127.0.0.1/regional_economy.industry") \  
    .load()  
  
education_df = spark.read.format("mongo") \  
    .option("uri", "mongodb://127.0.0.1/regional_economy.education") \  
    .load()  
  
# Cache the finance DataFrame if used multiple times  
finance_df.cache()
```

## Optimizing Performance

---

### Creating Indexes in MongoDB

After inserting the CSV data into MongoDB, we create indexes on frequently used fields (e.g., Year):

```
finance_collection.create_index([("Year", 1)])  
industry_collection.create_index([("Year", 1)])  
education_collection.create_index([("Year", 1)])
```

Since queries or aggregations commonly group, filter, or sort by the Year field, creating an index significantly reduces the amount of data to scan, improving lookup efficiency in MongoDB.

## Caching frequently used DataFrames

```
# Cache the finance dataframe to keep it in memory  
finance_df.cache()
```

For a DataFrame that will be used multiple times, caching it in memory (or in serialized form) prevents repeated computation and file reads, improving query speed.

## Broadcast Join

When performing a join operation where one of the tables is significantly smaller (for example, `education_df`), we can use `broadcast(education_df)` to optimize the join:

```
from pyspark.sql.functions import broadcast  
  
joined_df = finance_df.join(broadcast(education_df), on="Year", how="inner")
```

Broadcasting the smaller dataset to all executors avoids a full shuffle, significantly reducing the cost of the join operation.

## Result and Insights

---

### Key Findings

- **Education:** Higher education levels were found to be linked with higher incomes and better job opportunities. Some regions with higher levels of education showed better financial outcomes.
- **Industry and Jobs:** It was observed that certain jobs were associated with specific education levels, showing the impact of education on job types.

### Performance Evaluation

The speed of query execution was tested, and performance improved after optimizations were made. The execution time was reduced by 30% through the use of fewer partitions and caching.

### Visualizations

Charts were created to visualize trends, such as the relationship between education and income, and how GDP has changed over time.

## Conclusion and Future Work

---

### Challenges

- One challenge encountered was processing such a large dataset. It took some time to optimize the system for speed.
- Another difficulty was setting up the connection between MongoDB and Spark.

### Possible Improvements

- **Sharding:** The performance could be further improved by using **sharding** in MongoDB, which divides the data across multiple servers.
- **Machine Learning:** Future work could involve using **machine learning** to predict future job trends based on education and income data.

## Code Quality and Documentation

---

The code was organized and written clearly. A **README** file was included in the GitHub repository, explaining how to run the project and what each part of the code does.

## GitHub Repository

---

[GitHub Repository Link](#)